

## ▼ CNN

```

1 import numpy as np
2
3 import torch
4 import torchvision
5 import torchvision.transforms as transforms
6
7 import matplotlib.pyplot as plt
8 import os

```

```

1 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

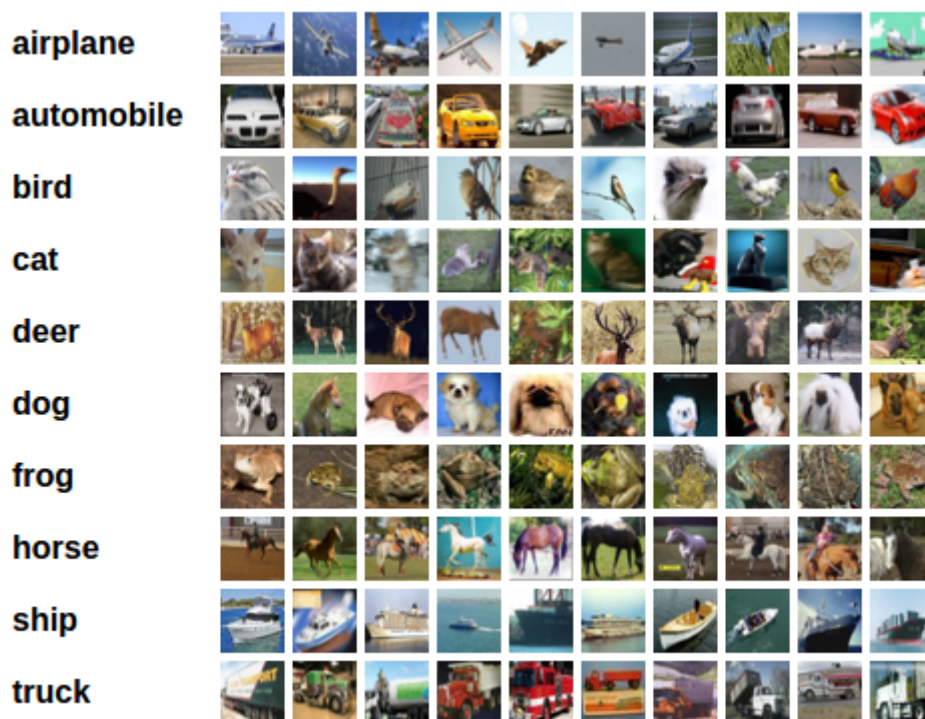
```

## ▼ CIFAR Image

```

1 from IPython.display import Image
2 image_url = 'https://bitbucket.org/hyuk125/lg_dic/raw/ae5f66c58a905dd778a803be'
3 Image(image_url)

```



## ▼ CIFAR data import

## ▼ OFFline pickle data

```

1 # def unpickle(file):
2 #     with open(file, 'rb') as fo:
3 #         dict = pickle.load(fo, encoding='bytes')
4 #     return dict
5 # data = unpickle(os.path.join(os.path.join(path, 'cifar-10-batches-py'), '/de

```

## ▼ OFFline PIL data

```

1 # OFFline PIL data
2 # trainset = torchvision.datasets.CIFAR10(root=path, train=True,
3 #                                         download=False)

```

## ▼ Online PIL image data

```

1 # # For online situation
2 trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
3                                         download=True)

```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to ./data/cifar-10-python  
170499072/? [00:05<00:00, 32370798.04it/s]

Extracting ./data/cifar-10-python.tar.gz to ./data

```

1 classes = ('plane', 'car', 'bird', 'cat',
2            'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

```

```


1 for i in range(5):
2     index = np.random.randint(len(trainset), dtype=int)
3     image, label = trainset[index]
4     image = image.resize((128, 128))
5     display(image)
6     print(classes[label])

```





plane



cat



ship

## ▼ Transform PIL data to torch data and normalization



```

1 transform = transforms.Compose(
2     [transforms.ToTensor(),
3      transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
4
5 trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
6                                         download=False, transform=transform)
7

```



```

1 a, _ =trainset[0]
2 a

```

```

tensor([[[[-0.5373, -0.6627, -0.6078, ..., 0.2392, 0.1922, 0.1608],
          [-0.8745, -1.0000, -0.8588, ..., -0.0353, -0.0667, -0.0431],
          [-0.8039, -0.8745, -0.6157, ..., -0.0745, -0.0588, -0.1451],
          ...,
          [ 0.6314, 0.5765, 0.5529, ..., 0.2549, -0.5608, -0.5843],
          [ 0.4118, 0.3569, 0.4588, ..., 0.4431, -0.2392, -0.3490],
          [ 0.3882, 0.3176, 0.4039, ..., 0.6941, 0.1843, -0.0353]],

        [[[-0.5137, -0.6392, -0.6235, ..., 0.0353, -0.0196, -0.0275],
          [-0.8431, -1.0000, -0.9373, ..., -0.3098, -0.3490, -0.3176],
          [-0.8118, -0.9451, -0.7882, ..., -0.3412, -0.3412, -0.4275],
          ...,
          [ 0.3333, 0.2000, 0.2627, ..., 0.0431, -0.7569, -0.7333],
          [ 0.0902, -0.0353, 0.1294, ..., 0.1608, -0.5137, -0.5843],
          [ 0.1294, 0.0118, 0.1137, ..., 0.4431, -0.0745, -0.2784]]],

        ...]])

```

```
[[-0.5059, -0.6471, -0.6627, ..., -0.1529, -0.2000, -0.1922],
 [-0.8431, -1.0000, -1.0000, ..., -0.5686, -0.6078, -0.5529],
 [-0.8353, -1.0000, -0.9373, ..., -0.6078, -0.6078, -0.6706],
 ...,
 [-0.2471, -0.7333, -0.7961, ..., -0.4510, -0.9451, -0.8431],
 [-0.2471, -0.6706, -0.7647, ..., -0.2627, -0.7333, -0.7333],
 [-0.0902, -0.2627, -0.3176, ..., 0.0980, -0.3412, -0.4353]]])
```

## ▼ Make Dataloader

```
1 trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
2                                           shuffle=True, num_workers=2)
```

## ▼ Define Model

```
1 import torch.nn as nn
2 import torch.nn.functional as F
```

### New modules

- `torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, ..., padding_mode='zeros', ...)`
- `torch.nn.MaxPool2d(kernel_size, stride=None, padding=0, ...)`

```
1 class Net(nn.Module):
2     def __init__(self):
3         super(Net, self).__init__()
4         self.conv1 = nn.Conv2d(3, 6, 5) # Convolution layer
5         self.pool = nn.MaxPool2d(2, 2) # Pooling layer
6
7         self.conv2 = nn.Conv2d(6, 16, 5)
8         self.fc1 = nn.Linear(16 * 5 * 5, 120)
9         self.fc2 = nn.Linear(120, 84)
10        self.fc3 = nn.Linear(84, 10)
11
12    def forward(self, x):
13        x = self.pool(F.relu(self.conv1(x)))
14        x = self.pool(F.relu(self.conv2(x)))
15        x = x.view(-1, 16 * 5 * 5)
16        x = F.relu(self.fc1(x))
17        x = F.relu(self.fc2(x))
18        x = self.fc3(x)
19        return x
20
21
22 net = Net()
23 net.to(device)
```

```

Net(
  (conv1): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
  (fc1): Linear(in_features=400, out_features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84, bias=True)
  (fc3): Linear(in_features=84, out_features=10, bias=True)
)

```

### Tips for complicated convolution layer (same as above model)

```

1 def conv_block(in_dim, out_dim):
2     model = nn.Sequential(
3         nn.Conv2d(in_dim, out_dim, kernel_size = 5),
4         nn.ReLU(),
5         nn.MaxPool2d(2, 2)
6     )
7     return model

```

```

1 class Net(nn.Module):
2     def __init__(self):
3         super(Net, self).__init__()
4         self.feature = nn.Sequential(
5             conv_block(3, 6),
6             conv_block(6, 16)
7         )
8
9         self.fc1 = nn.Linear(16 * 5 * 5, 120)
10        self.fc2 = nn.Linear(120, 84)
11        self.fc3 = nn.Linear(84, 10)
12
13    def forward(self, x):
14        # x = self.pool(F.relu(self.conv1(x)))
15
16        # x = self.conv1(x)
17        # x = self.pool(F.relu(self.conv2(x)))
18        x = self.feature(x)
19        x = x.view(-1, 16 * 5 * 5)
20        x = F.relu(self.fc1(x))
21        x = F.relu(self.fc2(x))
22        x = self.fc3(x)
23        return x
24
25
26 net = Net()
27 net.to(device)
28

```

## ▼ Learning the model

```

1 learning_rate = 0.001

1 import torch.optim as optim
2
3 criterion = nn.CrossEntropyLoss().to(device)
4 optimizer = optim.Adam(net.parameters(), lr = learning_rate)

1 epochs = 20

1 for epoch in range(epochs):
2     running_cost = 0.0
3
4     for step, (batch_data) in enumerate(trainloader):
5         batch_x, batch_y = batch_data[0].to(device), batch_data[1].to(device)
6
7         optimizer.zero_grad()
8
9         outputs = net(batch_x)
10        cost = criterion(outputs, batch_y)
11
12        cost.backward()
13        optimizer.step()
14
15        running_cost += cost.item()
16        if step % 2000 == 1999:
17            print('[%d, %5d] cost: %.3f' % (epoch + 1, step + 1, running_cost)
18            running_cost = 0.0
19

```

```

/usr/local/lib/python3.7/dist-packages/torch/nn/functional.py:718: UserWarning: Named tensors and all their associated APIs (nn.functional.max_pool2d, etc.) have been deprecated.
return torch.max_pool2d(input, kernel_size, stride, padding, dilation, ceil_mode)

```

```

[1, 2000] cost: 1.852
[1, 4000] cost: 1.612
[1, 6000] cost: 1.526
[1, 8000] cost: 1.456
[1, 10000] cost: 1.421
[1, 12000] cost: 1.374
[2, 2000] cost: 1.310
[2, 4000] cost: 1.304
[2, 6000] cost: 1.260
[2, 8000] cost: 1.252
[2, 10000] cost: 1.276
[2, 12000] cost: 1.237
[3, 2000] cost: 1.169
[3, 4000] cost: 1.180
[3, 6000] cost: 1.176
[3, 8000] cost: 1.171
[3, 10000] cost: 1.185
[3, 12000] cost: 1.176
[4, 2000] cost: 1.120
[4, 4000] cost: 1.104
[4, 6000] cost: 1.105

```

```

[4, 8000] cost: 1.127
[4, 10000] cost: 1.106
[4, 12000] cost: 1.102
[5, 2000] cost: 1.046
[5, 4000] cost: 1.041
[5, 6000] cost: 1.058
[5, 8000] cost: 1.082
[5, 10000] cost: 1.079
[5, 12000] cost: 1.080
[6, 2000] cost: 1.015
[6, 4000] cost: 1.013
[6, 6000] cost: 1.011
[6, 8000] cost: 1.034
[6, 10000] cost: 1.033
[6, 12000] cost: 1.053
[7, 2000] cost: 0.963
[7, 4000] cost: 0.968
[7, 6000] cost: 0.978
[7, 8000] cost: 1.006
[7, 10000] cost: 1.015
[7, 12000] cost: 0.993
[8, 2000] cost: 0.909
[8, 4000] cost: 0.962
[8, 6000] cost: 0.953
[8, 8000] cost: 0.963
[8, 10000] cost: 0.980
[8, 12000] cost: 0.990
[9, 2000] cost: 0.890
[9, 4000] cost: 0.931
[9, 6000] cost: 0.948
[9, 8000] cost: 0.949
[9, 10000] cost: 0.944
[9, 12000] cost: 0.965
[10, 2000] cost: 0.862
[10, 4000] cost: 0.904

```

## ▼ 정확도 판단

### ▼ Test dataset import

```

1 testset = torchvision.datasets.CIFAR10(root='./data', train=False,
2                                     download=True, transform=transform)
3 testloader = torch.utils.data.DataLoader(testset, batch_size=len(testset),
4                                     shuffle=False, num_workers=2)

```

Files already downloaded and verified

### ▼ Confusion matrix and scores

```

1 test_iter = iter(testloader)
2 test_x, test_labels = test_iter.next()

```

```

1 outputs = net(test_x.to(device))
2 _, predicted = torch.max(outputs, 1)

1 predicted

tensor([3, 8, 1, ..., 5, 1, 7], device='cuda:0')
```

## ▼ Confusion matrix

```

1 from sklearn.metrics import confusion_matrix
2 predicted = predicted.cpu()
3 print(confusion_matrix(test_labels, predicted))
```

```

[[691  44  31  35  34  15   8  33  55  54]
 [ 25 762   1  18   6   9  16  20  11 132]
 [ 87   7 383 103 147  82  74  83  15  19]
 [ 29  12  44 427  77 217  62  93  10  29]
 [ 21   4  42  96 541  59  91 130   8   8]
 [ 21   7  37 182  56 545  16 123   1  12]
 [  5   9  27 113  74  39 680  27  10  16]
 [ 10   2  13  63  73  79   8 735   2  15]
 [146  87  19  33  24  12   7  10 590  72]
 [ 43 103   5  26   8  22  16  62  15 700]]
```

## ▼ Precision

```

1 from sklearn.metrics import precision_score
2 print(precision_score(test_labels, predicted, average=None))
3 print(precision_score(test_labels, predicted, average='weighted'))
```

```

[0.64100186 0.73481196 0.63621262 0.38959854 0.52019231 0.50509731
 0.69529652 0.55851064 0.82287308 0.66225166]
0.6165846497338424
```

## ▼ Recall

```

1 from sklearn.metrics import recall_score
2 print(recall_score(test_labels, predicted, average=None))
3 print(recall_score(test_labels, predicted, average='weighted'))
```

```

[0.691 0.762 0.383 0.427 0.541 0.545 0.68  0.735 0.59  0.7  ]
0.6054
```

1